

MetroGraph — Strategy Dossier

A non-divergent family of artifacts auto-projected from the cross-domain Strategy OS corpus (*ingest render*). Every figure resolves to one `claim_id/source_id` and carries its evidence glyph; nothing is upgraded past its honest ceiling. Generated 2026-06-26.

MetroGraph — Investor Deck

Auto-generated by `ingest render` — a pure projection of `strategy.render_blocks`. Every figure carries its evidence glyph and resolves to one `claim_id / source_id`; the same block renders identically in every artifact (structural non-divergence). Do not hand-edit — re-run `ingest render`.

Headline metrics

Metric	Value	Evidence	Resolves to
Obtainable revenue (SOM), Monte-Carlo	\$37.6M / \$93.6M / \$226.3M (p10/p50/p90)	▮▮ modeled	market.claim.som
LTV:CAC (3-yr bounded), Monte-Carlo	2.6x / 8.5x / 28.2x (p10/p50/p90)	▮▮ modeled	finance.claim.ltv-cac-band
Comps valuation band	9.5x-37.4x revenue (median 17.5x, n=8)	○ speculative	finance.claim.valuation-band

Lead thesis: Lead wedge (schema-comprehension)

- supported-by-proxy / pending-experimental — ● supported-by-proxy
- cites* market.claim.graph-visualization-clutter-at-scale / voc.src.review-mining-2026q2

Risk: Canvas wedge erosion

- 15 dated rival moves converging on canvas-first editing — ○ speculative
- cites* compintel.claim.canvas-convergence / voc.src.review-mining-2026q2

Prioritized roadmap

Priority is derived (*cited-claim confidence × impact ÷ effort*) — it recomputes when facts change. Experiments are the needed-but-missing validation studies, emitted as first-class roadmap items.

Priority	Kind	Effort	Action
----------	------	--------	--------

0.350	 experiment	m	RUN THE MISSING STUDY to lift "force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/ schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).
0.169	positioning	m	Treat the canvas/spatial-editing wedge as CONTESTED, not open: 15 dated rival moves (n8n Canvas UI, Retool, Node-RED 5.0, ToolJet) are converging on canvas-first editing. Lead on schema-comprehension depth, not canvas novelty.
0.160	integration	m	Build the Postgres information_schema introspection integration first (derived priority 4.51, highest reach x fit / effort), then DuckDB and Snowflake.
0.116	 experiment	m	RUN THE MISSING STUDY to lift "graph-visualization-clutter-at-scale" past its proxy ceiling: Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV = time-on-task, error count, SUS; n>=30. Lands in voc.usability_sessions (currently EMPTY).
0.113	build	l	Ship the high-pain wedge features that the product corpus shows are still planned/in-progress (the differentiation pivot column) before the proxy-supported wedge can be tested in market.
0.087	 experiment	m	RUN THE MISSING STUDY to lift "direct-manipulation-outperforms-conversation-graph-exploration" past its proxy ceiling: Pre-registered preference + task-success test: visual graph view vs chat/conversational control of the same agent workflow; DV = success rate, trust calibration, time-to-detect-error; n>=40. Lands in voc.ab_experiments (currently EMPTY).
0.087	 experiment	m	RUN THE MISSING STUDY to lift "information-foraging-predicts-metro-map-adoption" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).
0.072	compliance	xl	Ship the enterprise compliance baseline (SSO/SAML, RBAC, MFA, immutable audit logging) before pursuing the enterprise-data-teams segment — 14 blocker gaps currently gate it.

Wedge evidence status (honest ceilings)

*The differentiating claims, re-graded against cross-domain evidence. The ceiling is **supported-by-proxy** — never experimental — because no MetroGraph behavioral study exists yet. Each names the study that would lift it.*

Wedge claim	Corpus verdict	Cross-domain ceiling	Needed study
agent-observability-through-visualization-improves-trust	disputed	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
mixed-initiative-design-ai-ui-parity-prevents-transparency-backfire	disputed	⚠ contested	Pre-registered controlled A/B: MetroGraph metro/ schematic layout vs force-directed on a sc...
mixed-initiative-requires-visualization-to-prevent-agent-opacity	refuted	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
visual-affordances-enable-interaction-without-training	refuted	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
direct-manipulation-outperforms-conversation-graph-exploration	refuted	🟡 supported-by-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem	supported	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/ schematic layout vs force-directed on a sc...
graph-visualization-clutter-at-scale	disputed	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...

information-foraging-predicts-metro-map-adoption	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
progressive-disclosure-unlocks-schema-acquisition-in-graphs	refuted	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...
wayfinding-in-schematic-maps-transfers-from-transit-knowledge	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
ai-ui-parity-exclusive-wedge	supported	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
flight-to-chat-when-ui-confuses-documented	supported	🟡 weak-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
schematic-maps-outperform-force-directed-database-exploration	refuted	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...

Red-team (standing bear case)

- 🟡 **FIRED** · *fatal* · **thesis-falsifier** — FALSIFIER: HCI literature is split and shows Sugiyama hierarchical can beat orthogonal; metro-layout superiority is only weak-proxy.
- Mitigation*: Run the pre-registered layout A/B (rec experiment-*) before betting positioning on layout superiority.
- 🟡 **FIRED** · *major* · **thesis-falsifier** — FALSIFIER: if incumbents ship canvas-first spatial editing, the canvas wedge commoditizes. Watch compintel for dated canvas moves.
- Mitigation*: Pivot the wedge narrative from canvas novelty to schema-comprehension depth + DB-introspection (rec defend-canvas-wedge).
- 🟡 **FIRED** · *major* · **competitor-countermove** — COUNTERMOVE: n8n/Zapier/Retool/Airtable are shipping first-class agent nodes; the agent-visibility wedge is contested.
- Mitigation*: Differentiate on agent-state VISUALIZATION depth, not agent execution (which is commoditizing fast).

All claims trace to the cross-domain corpus; proxy/contested grades are surfaced, not hidden. Generated from the committed snapshot — regenerate with `ingest render`.

MetroGraph — Strategy Memo

Auto-generated by `ingest render` — a pure projection of `strategy.render_blocks`. Every figure carries its evidence glyph and resolves to one `claim_id` / `source_id`; the same block renders identically in every artifact (structural non-divergence). Do not hand-edit — re-run `ingest render`.

Headline metrics

Metric	Value	Evidence	Resolves to
Obtainable revenue (SOM), Monte-Carlo	\$37.6M / \$93.6M / \$226.3M (p10/p50/p90)	📊 modeled	<code>market.claim.som</code>

Lead thesis: Lead wedge (schema-comprehension)

- supported-by-proxy / pending-experimental —  supported-by-proxy
- *cites* market.claim.graph-visualization-clutter-at-scale / voc.src.review-mining-2026q2

Risk: Enterprise gate

- 14 blocker compliance gaps (SSO/RBAC/MFA/audit/DPA/residency) —  supported
- *cites* governance.gap.no-sso-implementation / voc.src.review-mining-2026q2

Recommendation: First integration

- Postgres introspection (derived priority 4.51) —  supported-by-proxy
- *cites* ecosystem.integ.postgres-core / voc.src.review-mining-2026q2

Risk: Canvas wedge erosion

- 15 dated rival moves converging on canvas-first editing —  speculative
- *cites* compintel.claim.canvas-convergence / voc.src.review-mining-2026q2

Prioritized roadmap

Priority is derived (cited-claim confidence × impact ÷ effort) — it recomputes when facts change. Experiments are the needed-but-missing validation studies, emitted as first-class roadmap items.

Priority	Kind	Effort	Action
0.350	 experiment	m	RUN THE MISSING STUDY to lift "force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/ schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).
0.169	positioning	m	Treat the canvas/spatial-editing wedge as CONTESTED, not open: 15 dated rival moves (n8n Canvas UI, Retool, Node-RED 5.0, ToolJet) are converging on canvas-first editing. Lead on schema-comprehension depth, not canvas novelty.
0.160	integration	m	Build the Postgres information_schema introspection integration first (derived priority 4.51, highest reach x fit / effort), then DuckDB and Snowflake.
0.116	 experiment	m	RUN THE MISSING STUDY to lift "graph-visualization-clutter-at-scale" past its proxy ceiling: Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV = time-on-task, error count, SUS; n>=30. Lands in voc.usability_sessions (currently EMPTY).
0.113	build	l	Ship the high-pain wedge features that the product corpus shows are still planned/in-progress (the differentiation pivot column) before the proxy-supported wedge can be tested in market.
0.087	 experiment	m	RUN THE MISSING STUDY to lift "direct-manipulation-outperforms-conversation-graph-exploration" past its proxy ceiling: Pre-registered preference + task-success test: visual graph view vs chat/conversational control of the same agent workflow; DV = success rate, trust calibration, time-to-detect-error; n>=40. Lands in voc.ab_experiments (currently EMPTY).
0.087	 experiment	m	RUN THE MISSING STUDY to lift "information-foraging-predicts-metro-map-adoption" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).

0.072	compliance	xl	Ship the enterprise compliance baseline (SSO/SAML, RBAC, MFA, immutable audit logging) before pursuing the enterprise-data-teams segment — 14 blocker gaps currently gate it.
-------	------------	----	---

Wedge evidence status (honest ceilings)

The differentiating claims, re-graded against cross-domain evidence. The ceiling is **supported-by-proxy** — never experimental — because no MetroGraph behavioral study exists yet. Each names the study that would lift it.

Wedge claim	Corpus verdict	Cross-domain ceiling	Needed study
agent-observability-through-visualization-improves-trust	disputed	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
mixed-initiative-design-ai-ui-parity-prevents-transparency-backfire	disputed	⚠ contested	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
mixed-initiative-requires-visualization-to-prevent-agent-opacity	refuted	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
visual-affordances-enable-interaction-without-training	refuted	⚠ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
direct-manipulation-outperforms-conversation-graph-exploration	refuted	🟡 supported-by-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem	supported	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
graph-visualization-clutter-at-scale	disputed	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...
information-foraging-predicts-metro-map-adoption	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
progressive-disclosure-unlocks-schema-acquisition-in-graphs	refuted	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...
wayfinding-in-schematic-maps-transfers-from-transit-knowledge	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
ai-ui-parity-exclusive-wedge	supported	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
flight-to-chat-when-ui-confuses-documented	supported	🟡 weak-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
schematic-maps-outperform-force-directed-database-exploration	refuted	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...

Red-team (standing bear case)

- 🟡 **FIRED** · *fatal* · **thesis-falsifier** — FALSIFIER: HCI literature is split and shows Sugiyama hierarchical can beat orthogonal; metro-layout superiority is only weak-proxy.
- Mitigation*: Run the pre-registered layout A/B (rec experiment-*) before betting positioning on layout superiority.
- 🟡 **FIRED** · *major* · **thesis-falsifier** — FALSIFIER: if incumbents ship canvas-first spatial editing, the canvas wedge commoditizes. Watch compintel for dated canvas moves.
- Mitigation*: Pivot the wedge narrative from canvas novelty to schema-comprehension depth + DB-introspection (rec defend-canvas-wedge).

- **FIRED** · *major* · **competitor-countermove** — COUNTERMOVE: n8n/Zapier/Retool/Airtable are shipping first-class agent nodes; the agent-visibility wedge is contested.
- *Mitigation*: Differentiate on agent-state VISUALIZATION depth, not agent execution (which is commoditizing fast).

All claims trace to the cross-domain corpus; proxy/contested grades are surfaced, not hidden. Generated from the committed snapshot — regenerate with ingest render.

MetroGraph — Battlecard vs Neo4J

Auto-generated by `ingest render` — a pure projection of `strategy.render_blocks`. Every figure carries its evidence glyph and resolves to one `claim_id / source_id`; the same block renders identically in every artifact (structural non-divergence). Do not hand-edit — re-run `ingest render`.

Lead thesis: Lead wedge (schema-comprehension)

- supported-by-proxy / pending-experimental — **supported-by-proxy**
- *cites* `market.claim.graph-visualization-clutter-at-scale / voc.src.review-mining-2026q2`

Risk: Canvas wedge erosion

- 15 dated rival moves converging on canvas-first editing — **speculative**
- *cites* `compintel.claim.canvas-convergence / voc.src.review-mining-2026q2`

Recommendation: First integration

- Postgres introspection (derived priority 4.51) — **supported-by-proxy**
- *cites* `ecosystem.integ.postgres-core / voc.src.review-mining-2026q2`

All claims trace to the cross-domain corpus; proxy/contested grades are surfaced, not hidden. Generated from the committed snapshot — regenerate with ingest render.

MetroGraph — Board Update

Auto-generated by `ingest render` — a pure projection of `strategy.render_blocks`. Every figure carries its evidence glyph and resolves to one `claim_id / source_id`; the same block renders identically in every artifact (structural non-divergence). Do not hand-edit — re-run `ingest render`.

Headline metrics

Metric	Value	Evidence	Resolves to
Obtainable revenue (SOM), Monte-Carlo	\$37.6M / \$93.6M / \$226.3M (p10/p50/p90)	 modeled	market.claim.som
LTV:CAC (3-yr bounded), Monte-Carlo	2.6x / 8.5x / 28.2x (p10/p50/p90)	 modeled	finance.claim.ltv-cac-band

Risk: Enterprise gate

- 14 blocker compliance gaps (SSO/RBAC/MFA/audit/DPA/residency) — ✓ supported
- *cites* governance.gap.no-sso-implementation / voc.src.review-mining-2026q2

Prioritized roadmap

Priority is derived (*cited-claim confidence* × *impact* ÷ *effort*) — it recomputes when facts change. Experiments are the needed-but-missing validation studies, emitted as first-class roadmap items.

Priority	Kind	Effort	Action
0.350	 experiment	m	RUN THE MISSING STUDY to lift "force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).
0.169	positioning	m	Treat the canvas/spatial-editing wedge as CONTESTED, not open: 15 dated rival moves (n8n Canvas UI, Retool, Node-RED 5.0, ToolJet) are converging on canvas-first editing. Lead on schema-comprehension depth, not canvas novelty.
0.160	integration	m	Build the Postgres information_schema introspection integration first (derived priority 4.51, highest reach x fit / effort), then DuckDB and Snowflake.
0.116	 experiment	m	RUN THE MISSING STUDY to lift "graph-visualization-clutter-at-scale" past its proxy ceiling: Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV = time-on-task, error count, SUS; n>=30. Lands in voc.usability_sessions (currently EMPTY).
0.113	build	l	Ship the high-pain wedge features that the product corpus shows are still planned/in-progress (the differentiation pivot column) before the proxy-supported wedge can be tested in market.
0.087	 experiment	m	RUN THE MISSING STUDY to lift "direct-manipulation-outperforms-conversation-graph-exploration" past its proxy ceiling: Pre-registered preference + task-success test: visual graph view vs chat/conversational control of the same agent workflow; DV = success rate, trust calibration, time-to-detect-error; n>=40. Lands in voc.ab_experiments (currently EMPTY).
0.087	 experiment	m	RUN THE MISSING STUDY to lift "information-foraging-predicts-metro-map-adoption" past its proxy ceiling: Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a schema-comprehension + path-finding task; DV = task time & error rate; n>=40 within-subject. Lands in voc.ab_experiments / hci primary_studies (currently EMPTY).
0.072	compliance	xl	Ship the enterprise compliance baseline (SSO/SAML, RBAC, MFA, immutable audit logging) before pursuing the enterprise-data-teams segment — 14 blocker gaps currently gate it.

Wedge evidence status (honest ceilings)

The differentiating claims, re-graded against cross-domain evidence. The ceiling is **supported-by-proxy** — never experimental — because no MetroGraph behavioral study exists yet. Each names the study that would lift it.

Wedge claim	Corpus verdict	Cross-domain ceiling	Needed study
agent-observability-through-visualization-improves-trust	disputed	⚠️ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
mixed-initiative-design-ai-ui-parity-prevents-transparency-backfire	disputed	⚠️ contested	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
mixed-initiative-requires-visualization-to-prevent-agent-opacity	refuted	⚠️ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
visual-affordances-enable-interaction-without-training	refuted	⚠️ contested	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
direct-manipulation-outperforms-conversation-graph-exploration	refuted	🟡 supported-by-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
force-directed-graph-layout-remains-dominant-but-unoptimized-for-schem	supported	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
graph-visualization-clutter-at-scale	disputed	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...
information-foraging-predicts-metro-map-adoption	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
progressive-disclosure-unlocks-schema-acquisition-in-graphs	refuted	🟡 supported-by-proxy	Usability study: surface-area / progressive-disclosure variant vs full-canvas control; DV ...
wayfinding-in-schematic-maps-transfers-from-transit-knowledge	refuted	🟡 supported-by-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
ai-ui-parity-exclusive-wedge	supported	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...
flight-to-chat-when-ui-confuses-documented	supported	🟡 weak-proxy	Pre-registered preference + task-success test: visual graph view vs chat/conversational co...
schematic-maps-outperform-force-directed-database-exploration	refuted	🟡 weak-proxy	Pre-registered controlled A/B: MetroGraph metro/schematic layout vs force-directed on a sc...

Red-team (standing bear case)

- **FIRED** · *fatal* · **thesis-falsifier** — FALSIFIER: HCI literature is split and shows Sugiyama hierarchical can beat orthogonal; metro-layout superiority is only weak-proxy.
- *Mitigation*: Run the pre-registered layout A/B (rec experiment-*) before betting positioning on layout superiority.
- **FIRED** · *major* · **thesis-falsifier** — FALSIFIER: if incumbents ship canvas-first spatial editing, the canvas wedge commoditizes. Watch compintel for dated canvas moves.
- *Mitigation*: Pivot the wedge narrative from canvas novelty to schema-comprehension depth + DB-introspection (rec defend-canvas-wedge).
- **FIRED** · *major* · **competitor-countermove** — COUNTERMOVE: n8n/Zapier/Retool/Airtable are shipping first-class agent nodes; the agent-visibility wedge is contested.
- *Mitigation*: Differentiate on agent-state VISUALIZATION depth, not agent execution (which is commoditizing fast).

All claims trace to the cross-domain corpus; proxy/contested grades are surfaced, not hidden. Generated from the committed snapshot — regenerate with `ingest render`.

